

Importing Required Libraries

In this step, necessary Python libraries are imported for data handling, text preprocessing, feature extraction, model training, and evaluation.

```
In [1]: import pandas as pd
import numpy as np
import re

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_mat

import matplotlib.pyplot as plt
import seaborn as sns
import pickle
```

Dataset Loading

The fake news dataset is loaded using the Pandas library. The dataset contains news statements and corresponding labels indicating whether the news is fake or genuine.

```
In [2]: # Load dataset (update path if needed)
data = pd.read_csv("train.csv")

# Display first 5 rows
data.head()
```

Out[2]:

| | | Statement | Label |
|---|---|-----------|-------|
| 0 | Says the Annies List political group supports ... | | False |
| 1 | When did the decline of coal start? It started... | | True |
| 2 | Hillary Clinton agrees with John McCain "by vo... | | True |
| 3 | Health care reform legislation is likely to ma... | | False |
| 4 | The economic turnaround started at the end of ... | | True |

```
In [4]: # Dataset information
data.info()

# Label distribution
data['Label'].value_counts()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10240 entries, 0 to 10239
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype  
---  --          -----          --    
 0   Statement    10240 non-null   object 
 1   Label        10240 non-null   bool    
dtypes: bool(1), object(1)
memory usage: 90.1+ KB

Out[4]: Label
True      5752
False     4488
Name: count, dtype: int64
```

Text Preprocessing

This step cleans the news text by removing special characters, converting text to lowercase, and eliminating unnecessary spaces. Text preprocessing improves model accuracy.

```
In [5]: def preprocess_text(text):
    text = re.sub(r'[^a-zA-Z]', ' ', str(text))
    text = text.lower()
    text = text.split()
    text = ' '.join(text)
    return text
```

```
In [7]: data['clean_text'] = data['Statement'].apply(preprocess_text)
data[['Statement', 'clean_text']].head()
```

| | Statement | clean_text |
|---|---|---|
| 0 | Says the Annies List political group supports | says the annies list political group supports |
| 1 | When did the decline of coal start? It started... | when did the decline of coal start it started |
| 2 | Hillary Clinton agrees with John McCain "by vo... | hillary clinton agrees with john mccain by vot... |
| 3 | Health care reform legislation is likely to ma... | health care reform legislation is likely to ma... |
| 4 | The economic turnaround started at the end of ... | the economic turnaround started at the end of ... |

Feature Extraction using TF-IDF

TF-IDF (Term Frequency–Inverse Document Frequency) is used to convert textual data into numerical feature vectors based on word importance.

```
In [9]: tfidf = TfidfVectorizer(stop_words='english', max_df=0.7)

X = tfidf.fit_transform(data['clean_text'])
```

```
y = data['Label']

print("Feature matrix shape:", X.shape)

Feature matrix shape: (10240, 11318)
```

Model Training

A Logistic Regression model is trained using TF-IDF features to classify news articles as fake or genuine.

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [11]: model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

```
Out[11]: ▾ LogisticRegression ⓘ ?
```

► Parameters

Model Evaluation

The trained model is evaluated using accuracy, precision, recall, F1-score, and confusion matrix to measure classification performance.

```
In [12]: y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

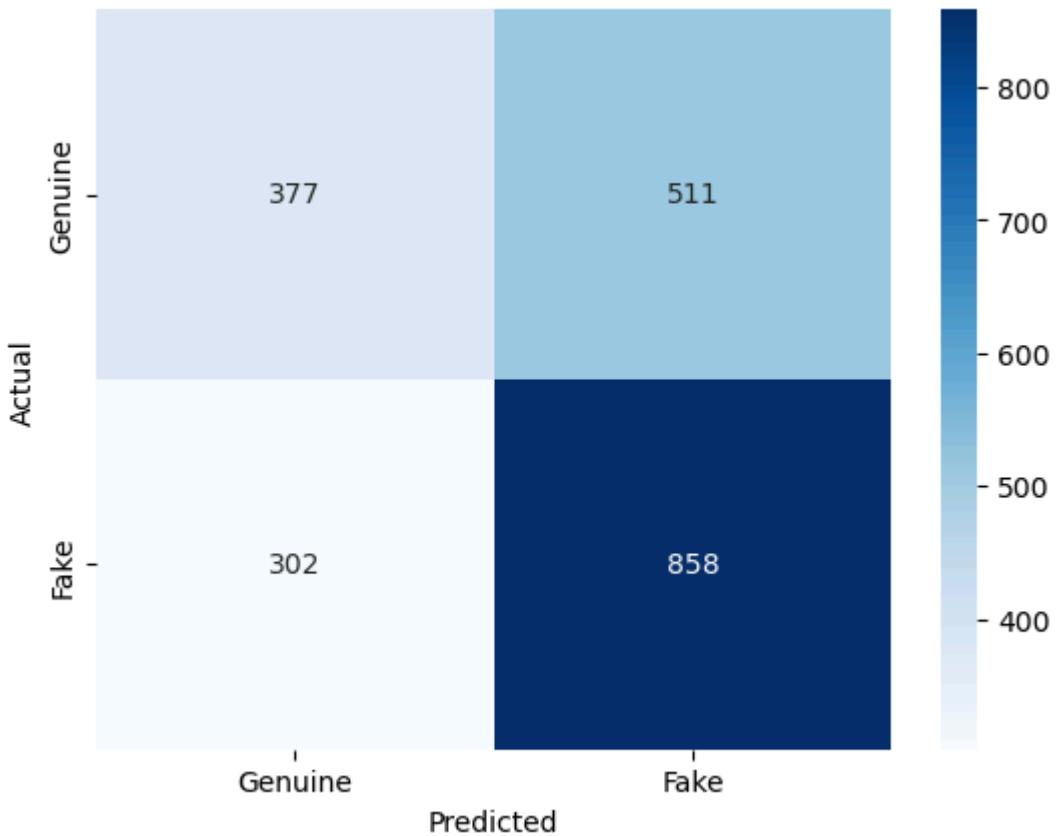
Accuracy: 0.60302734375

| Classification Report: | | | | |
|------------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| False | 0.56 | 0.42 | 0.48 | 888 |
| True | 0.63 | 0.74 | 0.68 | 1160 |
| accuracy | | | 0.60 | 2048 |
| macro avg | 0.59 | 0.58 | 0.58 | 2048 |
| weighted avg | 0.60 | 0.60 | 0.59 | 2048 |

```
In [13]: cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Genuine', 'Fake'],
            yticklabels=['Genuine', 'Fake'])

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



Prediction on New Data

The trained model predicts whether a given news statement is fake or genuine based on learned patterns.

```
In [14]: def predict_news(news_text):
    news_text = preprocess_text(news_text)
    vector = tfidf.transform([news_text])
    prediction = model.predict(vector)
    return "Fake News" if prediction[0] == 1 else "Genuine News"
```

```
In [15]: sample_news = "Breaking news: Scientists discover a cure for cancer!"
print(predict_news(sample_news))
```

Genuine News

Saving the Model

The trained model and TF-IDF vectorizer are saved for future use without retraining.

```
In [16]: with open("final_model.sav", "wb") as model_file:
    pickle.dump(model, model_file)

    with open("tfidf_vectorizer.sav", "wb") as vectorizer_file:
        pickle.dump(tfidf, vectorizer_file)

    print("Model and Vectorizer saved successfully!")
```

Model and Vectorizer saved successfully!

```
In [17]: with open("final_model.sav", "rb") as model_file:  
    loaded_model = pickle.load(model_file)  
  
    with open("tfidf_vectorizer.sav", "rb") as vectorizer_file:  
        loaded_tfidf = pickle.load(vectorizer_file)  
  
    test_text = "Government confirms new education policy reforms"  
    vector = loaded_tfidf.transform([preprocess_text(test_text)])  
    print("Prediction:", "Fake" if loaded_model.predict(vector)[0] == 1 else "Genuine")
```

Prediction: Fake